

# FOREX DATA PIPELINE USING APACHE AIRFLOW

Naveen Raju Sreerama Raju Govinda Raju  
[naveenraju100@gmail.com](mailto:naveenraju100@gmail.com) | +1 224 706 7718 | +91 9886157101 |  
<https://www.linkedin.com/in/naveen-raju-s-g-bb1486124>  
<https://naveenrajusg.github.io/Portfolio/>

## Table of Contents

- 1) Overview
- 2) Install Docker for Windows
- 3) Install Airflow steps
- 4) Docker
- 5) Flow chart of Forex Data pipeline
- 6) Architecture back end details
- 7) Pipeline code
  - A) Create a Gist in our Github account
  - B) Define DAG
  - C) Create a task named “is\_forex\_rates\_available” - HttpSensor to check if the API is available and start a docker container airflow airflow-section-3
  - D) Create task named “if currency file is available” - File sensor and check
  - E) Create task named “downloading\_rates” to download the forex rates from API - Python operator and check
  - F) Create a task named “forex\_rates” - Bash operator to save the forex rates into HDFS and check it
  - G) Create a task named “creating\_forex\_rates\_table” - Hive operator and check it
  - H) Process the forex rates with Spark - Spark submit operator
  - I) Send Email notifications - Email Operator
  - J) Send Slack notifications - SlackWebHookOperator
  - K) Add dependency between tasks
  - L) Trigger a DAG from airflow UI:

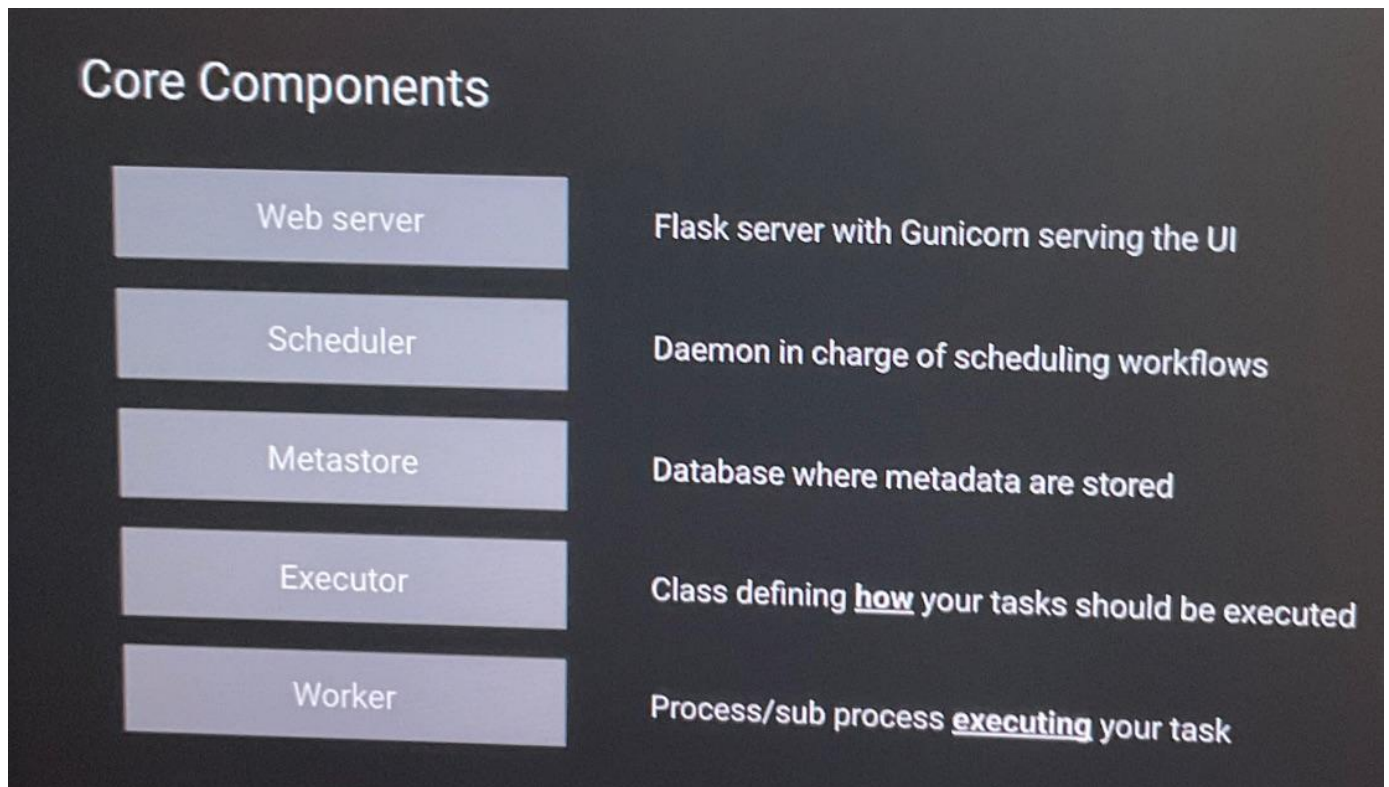
## 1) Overview

### Airflow :

Apache Airflow is an open-source platform used for orchestrating complex workflows and data pipelines.

Airflow allows you to define, schedule, and monitor workflows as Directed Acyclic Graphs (DAGs). Each DAG represents a workflow that consists of a set of tasks and their dependencies. Tasks can be anything from simple data processing steps to more complex operations, such as data extraction, transformation, loading, and model training.

### Core Components of Airflow :



### Directed Acyclic Graph:

"DAG" stands for "Directed Acyclic Graph." It is a fundamental concept used to represent workflows as a collection of tasks and their dependencies.

**Directed:** This means that there is a defined direction or flow between tasks.

**Acyclic:** A graph is acyclic if there are no cycles, which means there are no closed loops of dependencies.

A typical DAG represents a workflow with multiple tasks that need to be executed in a specific order, where the output of one task is often used as input for subsequent tasks. Each node in the DAG represents a task, and the edges between nodes represent the dependencies between tasks.

### Operator:

Operator is an object that encapsulates the task, the job that we want to execute. For example, if we want to connect to a database and insert data we use special operator to do that. An object encapsulates the tasks.

Categories of Operators : Action operator, Transfer operator, Sensor operator


### Task Instance:

When operator run on a dag then it is an instance.

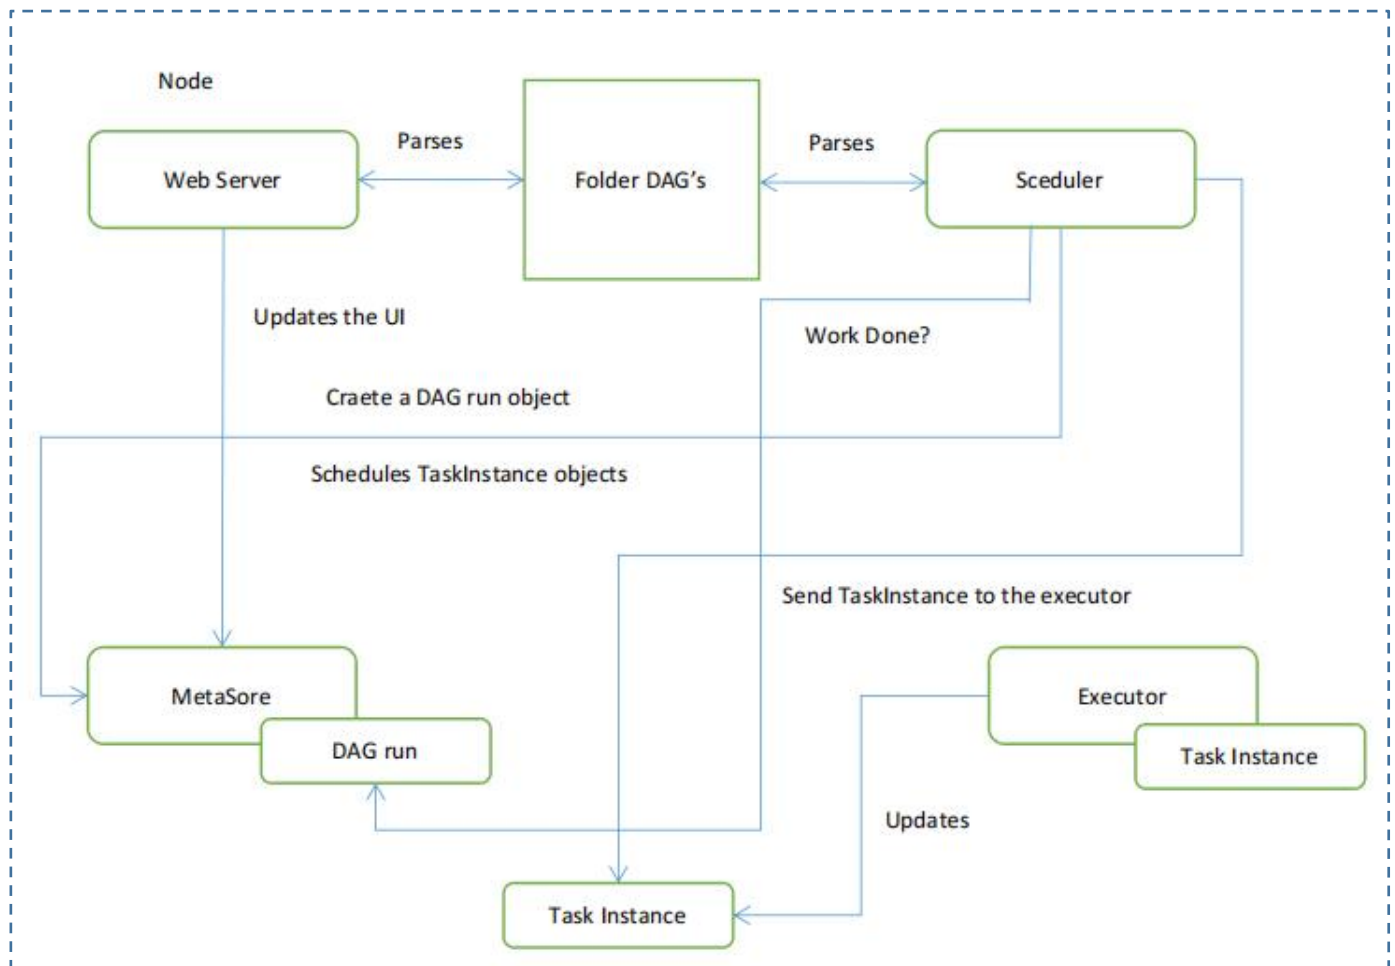
### One - node Architecture:

First the Web server fetches some metadata from the meta database in order to display information corresponding to Dag's, your task instances or your users on the user interface. Next, this category interacts with the meta database and the executor in order to trigger your Dag's, in order to trigger your tasks.

Finally, the executor interacts also with the meta in order to update the tasks that just have been completed. One thing to remember is that this scheduler, executor and the Web interact together with help of meta database. The executor has an internal queue and this is part of the executor. If you use the local executor, for example, and this is how your tasks are executed in the right order because there is a queue (Queue can be RabbitMQ or Redis) in the executor, that's why it executes tasks one after the other. By default we get sequential executor. And if we want to start setting up local executor where our tasks are executed in sub processes with both executors there is a queue in it and that's how your tasks are executed in order.



If we need to scale this to multi node then we can go with multi-node Celery executor or Kubernetes executor.



### Steps of DAG run in Apache Airflow:

- When new DAG is in DAG's folder both Web Server and Scheduler will parse the DAG.
- Scheduler checks if DAG is ready to be triggered if so DAG run object is created. (DAG run object is nothing but instance of our DAG running at a given time). This DAG run object is stored in MetaStore of Airflow with status running.
- If there is a task ready to be triggered in your DAG in that case scheduler creates TaskInstance objects corresponding to the task with the status scheduled in MetaStore of Airflow.
- Then Scheduler sends TaskInstance object to Executor with status queued.
- Once executor is ready to run the task this time TaskInstance object has status running and now Executor updates status of task in MetaStore.
- As soon as task is completed the Executor updates status of task in MetaStore. After this Scheduler verifies if there is no more task to run in DAG, if so DAG run object will now have status complete.
- Lastly Web Server updates UI.

## **2) Install Docker for Windows**

Follow the steps in the following link for installation <https://docs.docker.com/desktop/install/windows-install/>

## **3) Install Airflow steps**

Installing Airflow

```
docker run -it --rm -p 8080:8080 python:3.8-slim /bin/bash
```

\* Create and start a docker container from the Docker image python:3.8-slim and execute the command /bin/bash in order to have a shell session

`python -V`

\* Print the Python version

`export AIRFLOW_HOME=/usr/local/airflow`

\* Export the environment variable AIRFLOW\_HOME used by Airflow to store the dags folder, logs folder and configuration file

`env | grep airflow`

\* To check that the environment variable has been well exported

`apt-get update -y && apt-get install -y wget libczmq-dev curl libssl-dev git inetutils-telnet bind9utils freetds-dev libkrb5-dev libsasl2-dev libffi-dev libpq-dev freetds-bin build-essential default-libmysqlclient-dev apt-utils rsync zip unzip gcc && apt-get clean`

\* Install all tools and dependencies that can be required by Airflow

`useradd -ms /bin/bash -d ${AIRFLOW_HOME} airflow`

\* Create the user airflow, set its home directory to the value of AIRFLOW\_HOME and log into it

`cat /etc/passwd | grep airflow`

\* Show the file /etc/passwd to check that the airflow user has been created

`pip install --upgrade pip`

\* Upgrade pip (already installed since we use the Docker image python 3.5)

`su - airflow`

\* Log into airflow

`python -m venv .sandbox`

\* Create the virtual env named sandbox

`source .sandbox/bin/activate`

\* Activate the virtual environment sandbox

`wget https://raw.githubusercontent.com/apache/airflow/constraints-2.0.2/constraints-3.8.txt`

\* Download the requirement file to install the right version of Airflow's dependencies

`pip install "apache-airflow[crypto,celery,postgres,cncf.kubernetes,docker]"==2.0.2 --constraint ./constraints-3.8.txt`

\* Install the version 2.0.2 of apache-airflow with all subpackages defined between square brackets. (Notice that you can still add subpackages after all, you will use the same command with different subpackages even if Airflow is already installed)

airflow db init

\* Initialise the metadatabase

airflow scheduler &

\* Start Airflow's scheduler in background

airflow webserver &

\* Start Airflow's webserver in background

docker build -t airflow-basic .

\* Build a docker image from the Dockerfile in the current directory (airflow-materials/airflow-basic) and name it airflow-basic

docker image ls

\*to list all docker image

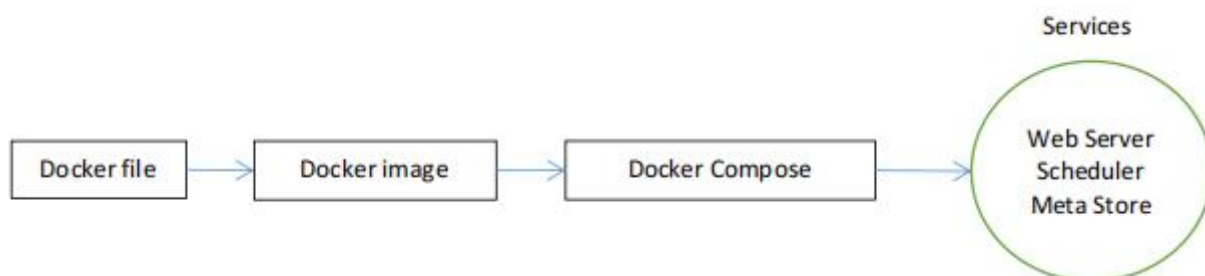
docker run --rm -d -p 8080:8080 airflow-basic

\*Airflow UI is running inside docker container, thus to access Airflow UI we should connect Docker container port 8080 to our machine port 8080

#### 4) Docker



Docker image is created from Docker file. Docker Image is like an application compiled. Docker file consists of all dependency installation commands.



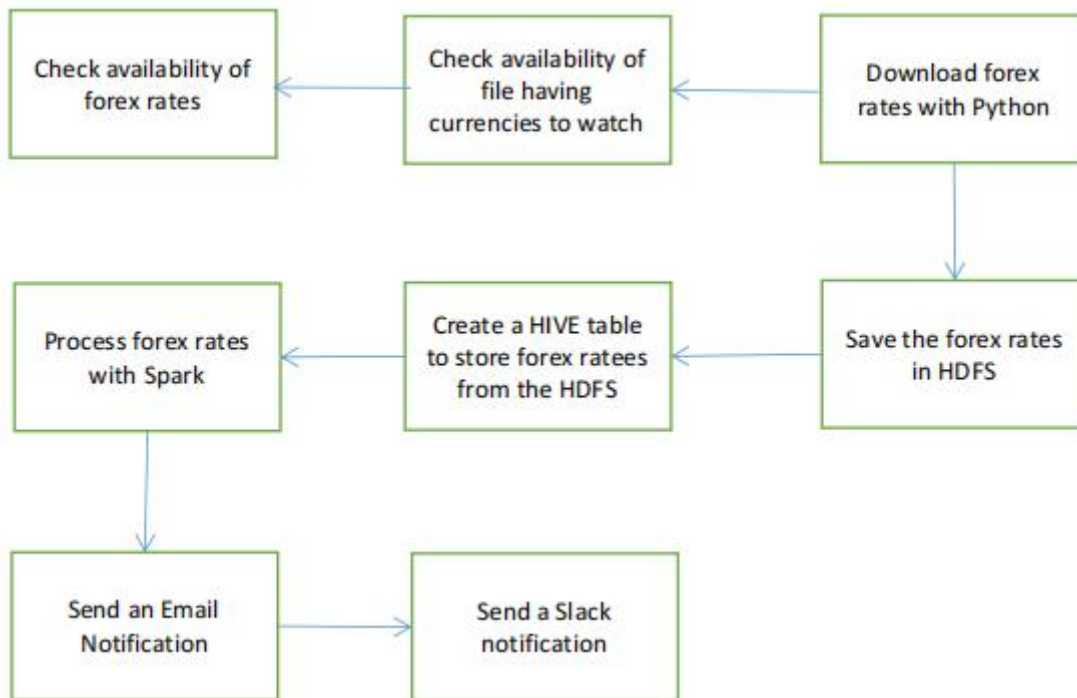
Docker Compose : is a tool provided by Docker that allows to define and manage multi containers. Networks and volumes as a single application, making it easier to orchestrate and deploy complex containerized applications.

In Airflow we have 3 main components MetaStore, Scheduler and Web server where each run in 3 different container managed by docker compose as single application.

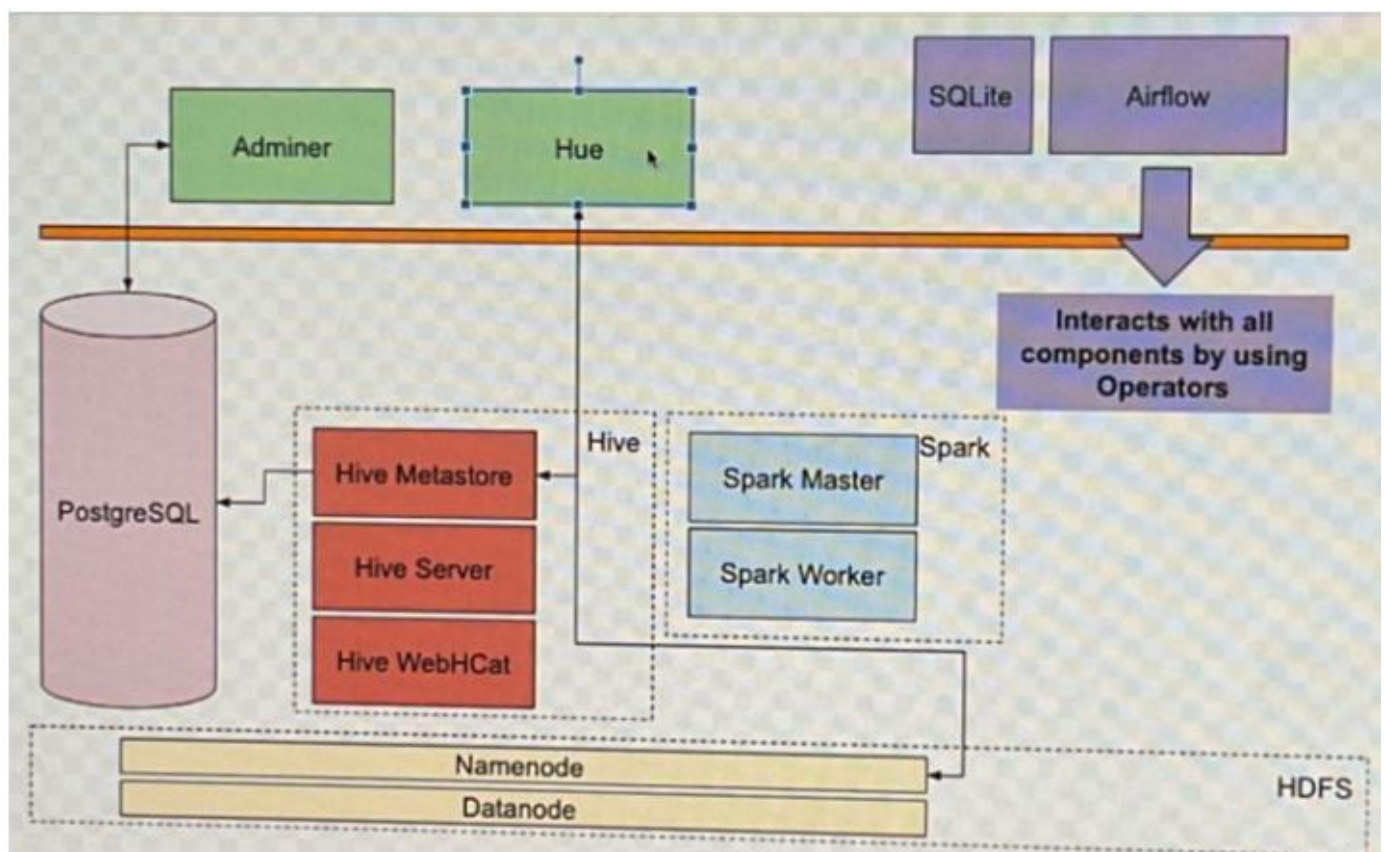


All 3 containers will run inside same network and so each container will be able to communicate with the other.

## 5) Flow chart of Forex Data pipeline



## 6) Architecture back end details



HDFS : This is where Forex data is stored

Admirer : is a tool to interact with PostgreSQL

Spark module : This consists of Spark Master and Spark Worker, is used to process Forex data

Hive module : Which consists of Hive Metastore, Hive server and Hive WebHCat, is used to query Forex data using SQL like statements, and also it uses PostgreSQL.

PostgreSQL :

PostgreSQL, often referred to as Postgres, is an open-source relational database management system (RDBMS) known for its robustness, reliability, and extensive feature set. It is a powerful and scalable database system that offers high performance, data integrity, and advanced capabilities for handling complex data.

PostgreSQL supports various operating systems, including Windows, macOS, Linux and BSD, making it a versatile choice for different environments. It adheres to SQL standards and provides support for a wide range of data types, indexing techniques, and advanced features such as triggers, stored procedures, and full-text search

Some notable features and capabilities of PostgreSQL include:

- **ACID Compliance:** PostgreSQL ensures Atomicity, Consistency, Isolation, and Durability, providing transactional integrity and reliability
- **Extensibility:** It supports user-defined data types, operators, functions, and extensions, allowing developers to customize and extend the database functionality.
- **Concurrency:** PostgreSQL employs multi version concurrency control (MVCC), enabling multiple transactions to access the database simultaneously without blocking each other
- **Replication and High Availability:** It offers various replication options, including and synchronous replication, for data redundancy and high availability.
- **Full-Text Search:** PostgreSQL provides robust full-text search capabilities, enabling efficient searching and indexing of textual data.
- **JSON and NoSQL Support:** It includes native support for JSON data, allowing storage, retrieval, and querying of JSON documents. It also provides support for NoSQL-like functionality
- **Scalability:** It supports scaling horizontally through sharding and can handle large volumes of data and high traffic workloads.

## **7) Pipeline code:**

### **A) Create a Gist in our Github account**

Create a Gist in our Github account with name 'api\_forex\_exchange.json' with 3 Json files in it.

```
api_forex_exchange.json
{
  "rates": {"CAD":1.31,"HKD":7.82,"ISK":121.32,"PHP":50.76,"DKK":6.73,"GBP":0.76,"JPY":108.56,"CHF":0.98,"EUR":0.90,"NZD":1.52,"USD":1.0,"SGD":1.35,"AUD":1.45},
  "base":"USD",
  "date":"2021-01-01"
```



```

}
api_forex_exchange_eur.json
{
  "rates": {"CAD":1.21,"GBP":0.36,"JPY":101.89,"USD":1.13,"NZD":1.41,"EUR":1.0},
  "base":"EUR",
  "date":"2021-01-01"
}

api_forex_exchange_usd.json
{
  "rates": {"CAD":1.31,"GBP":0.76,"JPY":108.56,"EUR":0.90,"NZD":1.52,"USD":1.0},
  "base":"USD",
  "date":"2021-01-01"
}

```

## **B) Define DAG**

```

from airflow import DAG
from datetime import datetime, timedelta

default_args = {
    "owner": "airflow",
    "email_on_failure": False,
    "email_on_retry": False,
    "email": "admin@localhost.com",
    "retries": 1, # On failure retry for one time
    "retry_delay": timedelta(minutes=5) # During the retry maximum processing time is 5minutes
}

with DAG("forex_data_pipeline", start_date=datetime(2021, 1, 1),
        schedule_interval="@daily", default_args=default_args, catchup=False) as dag:

# forex_data_pipeline - is name of DAG
# schedule_interval="@daily" - Execute DAG at midnight (00:00) every day
# catchup = False - Don't run DAG for older dates (I.e dates between current date and start date
mentioned)

```

## **C) Create a task named “is\_forex\_rates\_available” - HttpSensor to check if the API is available and start a docker container airflow airflow-section-3**

Write “is\_forex\_rates\_available” code in our DAG

```

is_forex_rates_available = HttpSensor(
    task_id="is_forex_rates_available",
    http_conn_id="forex_api",
    endpoint="naveenrajug/497e10579edfe65fdf1c3d60a387fa20",
    response_check=lambda response: "rates" in response.text,
    poke_interval=5,
    timeout=20
)

```

In current working directory of the pipeline code all command prompt code should be run:

```
In cmd:  
start.sh
```

```
start.sh
```

```
#!/bin/bash
```

```
# Build the base images from which are based the Dockerfiles  
# then Startup all the containers at once  
docker build -t hadoop-base docker/hadoop/hadoop-base && \  
docker build -t hive-base docker/hive/hive-base && \  
docker build -t spark-base docker/spark/spark-base && \  
docker-compose up -d --build
```

# `docker build -t hadoop-base docker/hadoop/hadoop-base`: This command builds a Docker image named "hadoop-base" using the Dockerfile located in the "docker/hadoop/hadoop-base" directory. This Docker image likely contains the base Hadoop installation.

# `docker build -t hive-base docker/hive/hive-base`: This command builds a Docker image named "hive-base" using the Dockerfile located in the "docker/hive/hive-base" directory. This Docker image likely contains the base Hive installation.

# `docker build -t spark-base docker/spark/spark-base`: This command builds a Docker image named "spark-base" using the Dockerfile located in the "docker/spark/spark-base" directory. This Docker image likely contains the base Spark installation.

# `docker-compose up -d --build`: This command brings up the Docker containers defined in the Docker Compose configuration file and rebuilds the containers if necessary. The `-d` flag runs the containers in detached mode, meaning they will run in the background. The `--build` flag ensures that any changes made to the Docker images are applied before starting the containers.

We can see a Docker container named `airflow-section-3` is created

The screenshot shows the Docker Desktop interface. On the left sidebar, the 'Containers' tab is selected. The main panel displays a table of containers. At the top, it shows 'Container CPU usage' as 285.61% / 1000% (10 cores allocated) and 'Container memory usage' as 2.81GB / 7.43GB. Below this is a search bar and a toggle for 'Only show running containers'. The table lists five containers: 'registry', 'vigilant\_bassi', 'confident\_chatelet', 'wizardly\_allen', and 'airflow-section-3'. The 'airflow-section-3' container is the only one in a 'Running' state, with a status of '(12/1)' and CPU usage of 285.61%. The other four containers are in an 'Exited' state. The bottom status bar shows 'RAM 7.84 GB', 'CPU 4.57%', and 'Not connected to Hub'.

	Name	Image	Status	CPU (%)	Port(s)	Last state	Actions
<input type="checkbox"/>	<a href="#">registry</a> cd48f0b2d105	<a href="#">registry:2</a>	Exited (2)	0%	5000:5000	9 hours	▶ ⋮ 🗑
<input type="checkbox"/>	<a href="#">vigilant_bassi</a> ebca03d9f651	<a href="#">python:3.7</a>	Exited (137)	0%		1 month	▶ ⋮ 🗑
<input type="checkbox"/>	<a href="#">confident_chatelet</a> 1ab0349fe954	<a href="#">python:3.7</a>	Exited (137)	0%		1 month	▶ ⋮ 🗑
<input type="checkbox"/>	<a href="#">wizardly_allen</a> d51b59051194	<a href="#">python:3.7</a>	Exited (137)	0%		1 month	▶ ⋮ 🗑
<input type="checkbox"/>	> <a href="#">airflow-section-3</a>		Running (12/1)	285.61%		5 minutes	▶ ⋮ 🗑

Showing 5 items

In browser type localhost:8080

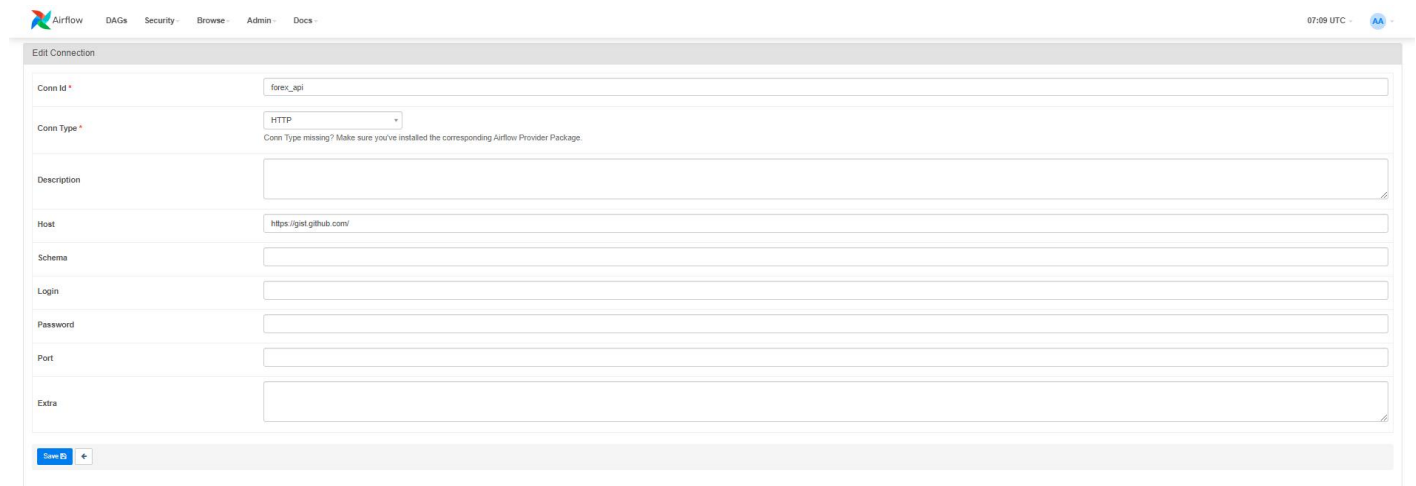
Username : airflow

Password : airflow

On Airflow UI navigate to Admin -> Connections -> +

Conn Id \* : forex\_api  
Conn Type \* : HTTP  
Host : <https://gist.github.com/>

Save



In cmd type : docker ps

Copy container id of airflow-section-3-airflow

```
E:\airflow course\airflow-materials\airflow-materials\airflow-section-3>docker ps
CONTAINER ID   IMAGE                                     COMMAND                  CREATED        STATUS        PORTS
TS
e8ccc6505d3ad  airflow-section-3-hive-webhcat          "./entrypoint.sh ./s..." 13 hours ago  Up 2 minutes (unhealthy) 100
00-10002/tcp, 50111/tcp
4224dd9ef67b   airflow-section-3-hue                  "./entrypoint.sh ./s..." 13 hours ago  Up 2 minutes (unhealthy) 0.0
.0.0:32762->8888/tcp
339bb2e4f3dc   airflow-section-3-livy                  "./entrypoint"           13 hours ago  Up 2 minutes (healthy) 0.0
.0.0:32758->8998/tcp
6f407b9843b2   airflow-section-3-hive-server           "./entrypoint.sh ./s..." 13 hours ago  Up 2 minutes (unhealthy) 100
01/tcp, 0.0.0.0:32760->10000/tcp, 0.0.0.0:32759->10002/tcp
dc7954b7e1f9   airflow-section-3-spark-worker          "./entrypoint.sh ./s..." 13 hours ago  Up 2 minutes (healthy) 100
00-10002/tcp, 0.0.0.0:32764->8081/tcp
d48b901f0aef   airflow-section-3-hive-metastore        "./entrypoint.sh ./s..." 13 hours ago  Up 2 minutes (unhealthy) 100
00-10002/tcp, 0.0.0.0:32761->9083/tcp
56adb4aaab4f   airflow-section-3-datanode              "./entrypoint.sh ./s..." 13 hours ago  Up 2 minutes (unhealthy) 986
4/tcp
eca31f0eebf2   airflow-section-3-spark-master          "./entrypoint.sh ./s..." 13 hours ago  Up 2 minutes (healthy) 606
6/tcp, 10000-10002/tcp, 0.0.0.0:32765->7077/tcp, 0.0.0.0:32766->8082/tcp
c9161445c8f6   wodby/adminer:latest                   "/entrypoint.sh php ..." 13 hours ago  Up 2 minutes (healthy) 0.0
.0.0:32767->9000/tcp
c0f0eb309232   airflow-section-3-namenode              "./entrypoint.sh ./s..." 13 hours ago  Up 2 minutes (unhealthy) 0.0
.0.0:32763->9870/tcp
061768b0c9c0   airflow-section-3-postgres              "docker-entrypoint.s..." 13 hours ago  Up 2 minutes (healthy) 0.0
.0.0:32769->5432/tcp
079d4ae10809   airflow-section-3-airflow               "./entrypoint.sh ./s..." 13 hours ago  Up 2 minutes (healthy) 0.0
.0.0:8080->8080/tcp, 10000-10002/tcp
airflow
```

In cmd type : docker exec -it 079d4ae10809 /bin/bash

This command is used to open a bash terminal to access airflow CLI running inside container

```
E:\airflow course\airflow-materials\airflow-materials\airflow-section-3>docker exec -it 079d4ae10809 /bin/bash
airflow@079d4ae10809:/$
```

Check if task “is\_forex\_rates\_available” of the DAG “forex\_data\_pipeline” is working

In cmd type: `airflow tasks test forex_data_pipeline is_forex_rates_available 2023-01-01`

```
[2023-08-02 07:21:26,893] {taskinstance.py:1115} INFO - Executing <Task(HttpSensor): is_forex_rates_available> on 2023-01-01T00:00:00+00:00
/usr/local/lib/python3.7/dist-packages/airflow/configuration.py:345 DeprecationWarning: The sensitive_variable_fields option in [admin] has been moved to the sensitive_var_conn_names option in [core] - the old setting has been used, but please update your config.
[2023-08-02 07:21:26,960] {taskinstance.py:1254} INFO - Exporting the following env vars:
AIRFLOW_CTX_DAG_EMAIL=admin@localhost.com
AIRFLOW_CTX_DAG_OWNER=airflow
AIRFLOW_CTX_DAG_ID=forex_data_pipeline
AIRFLOW_CTX_TASK_ID=is_forex_rates_available
AIRFLOW_CTX_EXECUTION_DATE=2023-01-01T00:00:00+00:00
[2023-08-02 07:21:26,960] {http.py:101} INFO - Poking: marclamberti/f45f872dea4dfd3eaa015a4a1af4b39b
[2023-08-02 07:21:26,963] {base.py:79} INFO - Using connection to: id: forex_api. Host: https://gist.github.com/, Port: None, Schema: , Login: , Password: None, extra: {}
[2023-08-02 07:21:26,964] {http.py:140} INFO - Sending 'GET' to url: https://gist.github.com/marclamberti/f45f872dea4dfd3eaa015a4a1af4b39b
[2023-08-02 07:21:27,643] {base.py:248} INFO - Success criteria met. Exiting.
[2023-08-02 07:21:27,645] {taskinstance.py:1219} INFO - Marking task as SUCCESS. dag_id=forex_data_pipeline, task_id=is_forex_rates_available, execution_date=20230101T000000, start_date=20230802T072126, end_date=20230802T072127
```

## D) Create task named “if currency file is available” - File sensor and check:

On Airflow UI navigate to Admin -> Connections -> +

Conn Id \* : forex\_path

Conn Type \* : file(path)

Extra : {“path”: “/opt/airflow/dags/files”} # this is where we will be looking if file exists.

Save

The screenshot shows the 'Add Connection' form in the Airflow UI. The 'Conn Id' field is filled with 'forex\_path'. The 'Conn Type' dropdown is set to 'File (path)'. Below this, there is a message: 'Conn Type missing? Make sure you've installed the corresponding Airflow Provider Package.' The 'Description' field is empty. The 'Host', 'Schema', 'Login', 'Password', and 'Port' fields are also empty. The 'Extra' field contains the JSON string '{\"path\": \"/opt/airflow/dags/files\"}'. At the bottom left, there is a 'Save' button.

Write File sensor code in the DAG

```

from airflow.sensors.filesystem import FileSensor

is_forex_currencies_file_available = FileSensor(
    task_id="is_forex_currencies_file_available",
    fs_conn_id="forex_path",
    filepath="forex_currencies.csv",
    poke_interval=5,
    timeout=20
)

```

‘poke\_interval’ specifies the time interval (in seconds) at which the FileSensor checks for the existence of the specified file on the filesystem.

‘timeout’ defines the maximum time (in seconds) the FileSensor will wait for the file to become available on the filesystem.

We may wonder where is the following path “/opt/airflow/dags/files”, this is inside a docker container where Airflow is running.

```

In cmd:
docker exec -it 079d4ae10809 /bin/bash
ls
cd /opt/airflow/dags/files
ls
pwd

```

```

E:\airflow course\airflow-materials\airflow-materials\airflow-section-3>docker exec -it 079d4ae10809 /bin/bash
airflow@079d4ae10809:/$ ls
bin  dev  etc  home  lib64  mnt  proc  run  srv  sys  usr
boot  entrypoint.sh  hadoop-data  lib  media  opt  root  sbin  start-airflow.sh  tmp  var
airflow@079d4ae10809:/$ cd /opt/airflow/dags/files
airflow@079d4ae10809:~/dags/files$ ls
forex_currencies.csv  forex_rates.json
airflow@079d4ae10809:~/dags/files$ pwd
/opt/airflow/dags/files
airflow@079d4ae10809:~/dags/files$

```

Any files we keep in our local file system in “mnt/airflow/dags” will be in airflow container “/opt/airflow/dags/”.

Test if the task “is\_forex\_currencies\_file\_available” is working:

```
In cmd type: airflow tasks test forex_data_pipeline is_forex_currencies_file_available 2023-01-01
```

```

[2023-08-02 07:47:36,324] {filesystem.py:55} INFO - Poking for file forex_currencies.csv
[2023-08-02 07:47:36,331] {filesystem.py:61} INFO - Found File forex_currencies.csv last modified: 20230609161332
[2023-08-02 07:47:36,331] {base.py:248} INFO - Success criteria met. Exiting.
[2023-08-02 07:47:36,333] {taskinstance.py:1219} INFO - Marking task as SUCCESS. dag_id=forex_data_pipeline, task_id=is_forex_currencies_file_available, execution_date=20230101T000000, start_date=20230802T074736, end_date=20230802T074736
airflow@079d4ae10809:~/dags/files$

```

## E) Create task named “downloading\_rates” to download the forex rates from API - Python operator and check

```

from airflow.operators.python_operator import PythonOperator
import csv
import requests
import json

```



```
def download_rates():
    BASE_URL =
"https://gist.githubusercontent.com/naveenrajug/497e10579edfe65fdf1c3d60a387fa20/raw/"
    ENDPOINTS = {
        'USD': 'api_forex_exchange_usd.json',
        'EUR': 'api_forex_exchange_eur.json'
    }
    with open('/opt/airflow/dags/files/forex_currencies.csv') as forex_currencies:
        reader = csv.DictReader(forex_currencies, delimiter=';')
        for idx, row in enumerate(reader):
            base = row['base']
            with_pairs = row['with_pairs'].split(' ')
            indata = requests.get(f'{BASE_URL}{ENDPOINTS[base]}').json()
            outdata = {'base': base, 'rates': {}, 'last_update': indata['date']}
            for pair in with_pairs:
                outdata['rates'][pair] = indata['rates'][pair]
            with open('/opt/airflow/dags/files/forex_rates.json', 'a') as outfile:
                json.dump(outdata, outfile)
                outfile.write('\n')
# Inside the DAG definition add the below code

downloading_rates = PythonOperator(
    task_id="downloading_rates",
    python_callable=download_rates
)
```

Test if the task “downloading\_rates” is working:

In cmd type: airflow tasks test forex\_data\_pipeline downloading\_rates 2023-01-01

```
[2023-08-02 08:33:35,732] {taskinstance.py:1254} INFO - Exporting the following env vars:
AIRFLOW_CTX_DAG_EMAIL=admin@localhost.com
AIRFLOW_CTX_DAG_OWNER=airflow
AIRFLOW_CTX_DAG_ID=forex_data_pipeline
AIRFLOW_CTX_TASK_ID=downloading_rates
AIRFLOW_CTX_EXECUTION_DATE=2023-01-01T00:00:00+00:00
[2023-08-02 08:33:36,835] {python.py:151} INFO - Done. Returned value was: None
[2023-08-02 08:33:36,840] {taskinstance.py:1219} INFO - Marking task as SUCCESS. dag_id=forex_data_pipeline, task_id=downloading_rates, execution_date=20230101T000000, start_date=20230802T083335, end_date=20230802T083336
airflow@079d4ae10809:~/dags/files$
```

Now we can see file named forex\_rates.json created in local file system

“/mnt/airflow/dags/files/forex\_rates.json“ the same will be created inside Airflow docker container.

File contents:

```
{"base": "EUR", "rates": {"USD": 1.13, "NZD": 1.41, "JPY": 101.89, "GBP": 0.36, "CAD": 1.21}, "last_update": "2021-01-01"}
{"base": "USD", "rates": {"EUR": 0.9, "NZD": 1.52, "JPY": 108.56, "GBP": 0.76, "CAD": 1.31}, "last_update": "2021-01-01"}
```

## **F) Create a task named “forex\_rates” - Bash operator to save the forex rates into HDFS and check it**

In real time if files are huge we need to save them in HDFS.

Bash operator is an operator that allows us to execute a bash command or a script as a task within an Airflow DAG. It is one of the core operator provided by Airflow.

Define bash operator in the DAG, here we are trying to save forex\_rates.json created in last step into HDFS.

### HUE :

Hue (Hadoop User Experience) is an open-source web-based interface that provides a graphical user interface (GUI) for interacting with Apache Hadoop and its ecosystem components. It is designed to simplify and enhance the user experience for working with Hadoop and related tools. Hue offers a wide range of features and capabilities that make it easier for users to interact with Hadoop clusters, perform data analysis, and develop workflows. Some key features of

Hue include:

- File Browser: Allows users to navigate and manage files stored in Hadoop Distributed File System (HDFS) or other compatible file systems.
- Query Editors: Provides interactive editors for writing and executing queries in languages like Hive, Impala, Pig, and Spark SQL it includes features like syntax highlighting, auto-completion, and result visualization.
- Job Designer: Enables users to visually design and schedule workflows using tools like Oozie and Apache Workflow Scheduler (AWS). It simplifies the creation and management of data pipelines.
- Data Browsing and Visualization Allows users to explore and analyze data stored in Hadoop using tools like Apache Hive, Apache Impala, and Apache Solr. It provides Interactive visualizations and data exploration capabilities.
- Security and User Management Offers features for managing user access, authentication, and authorization to Hadoop resources. It integrates with security mechanisms like and LDAP

Enables users to create and share dashboards for data visualization and reporting purposes. It supports various charting libraries and allows customization of dashboards

- Job Monitoring Provides monitoring and tracking capabilities for jobs running on the cluster. It allows users view job status, logs, and performance metrics.

```
from airflow.operators.bash_operator import BashOperator

saving_rates = BashOperator(
    task_id="saving_rates",
    bash_command="""
        hdfs dfs -mkdir -p /forex && \ # creates folder named forex in HDFS
        hdfs dfs -put -f $AIRFLOW_HOME/dags/files/forex_rates.json /forex    # Copy file from Airflow
        container to HDFS folder
    """
)
```

In browser type : <http://localhost:32762/> (It takes to HUE login)

Username : root

Password : root

Restart docker container if HUE port is not working

```
In cmd:
restart.sh
```

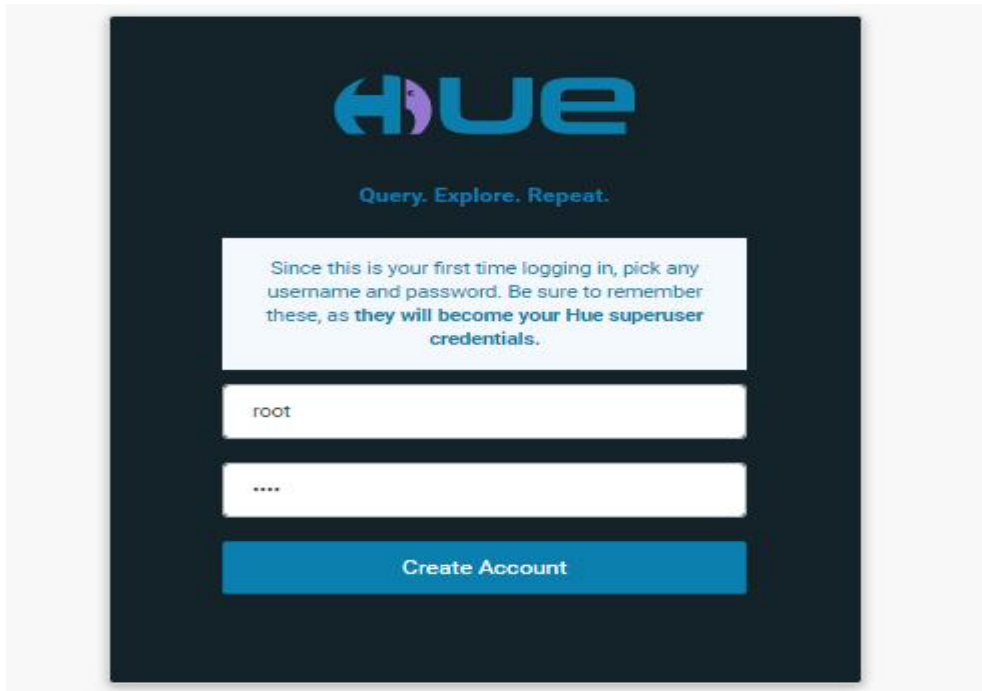


```
restart.sh
```

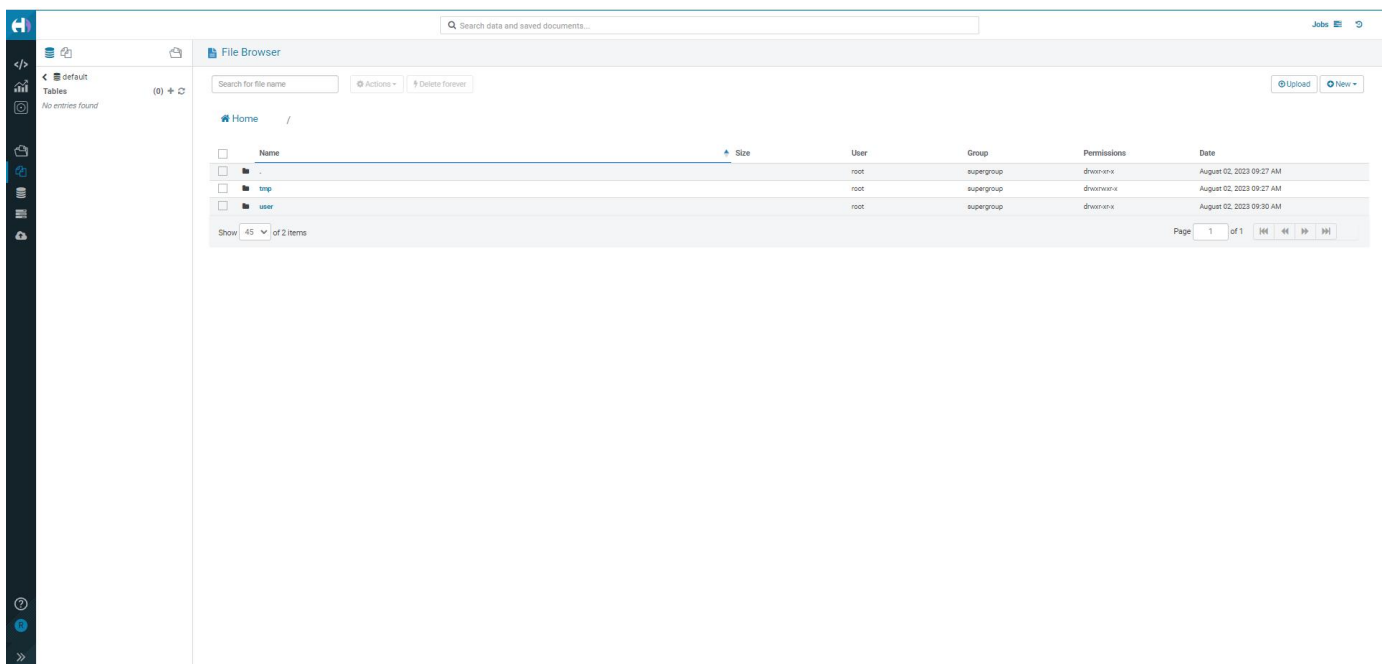
```
#!/bin/bash
```

```
./stop.sh
```

```
./start.sh
```



In HUE UI navigate to files to see files in HDFS



In cmd:

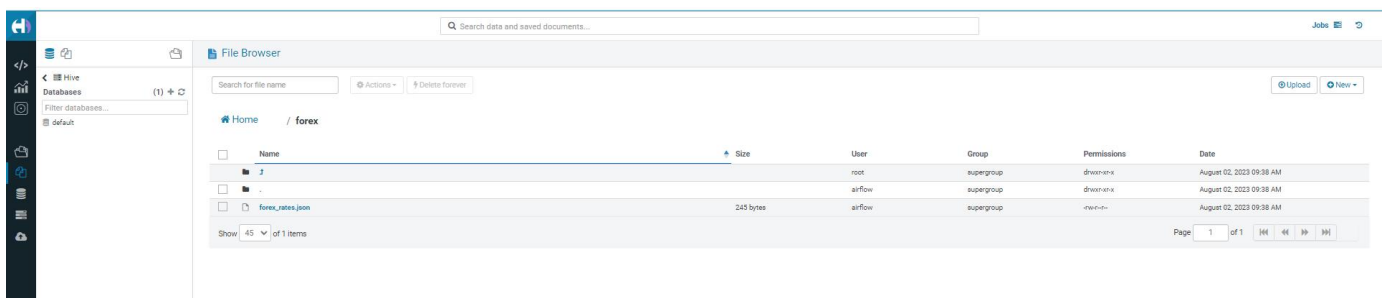
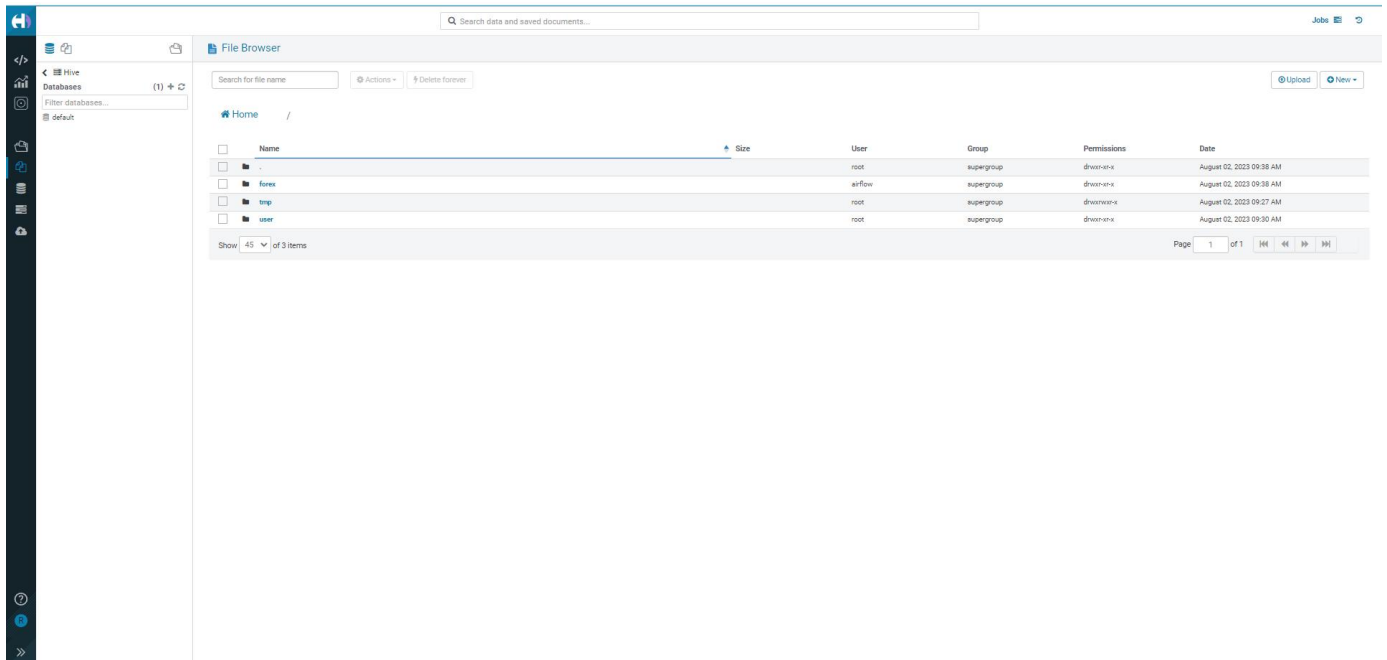
```
docker ps # copy container id of airflow-section-3-airflow
```

```
docker exec -it 510f580614fb /bin/bash
```

Test if the task “saving\_rates” is working:

In cmd type: `airflow tasks test forex_data_pipeline saving_rates 2023-01-01`

Now refresh the HUE UI, now we can see forex folder created with a file `forex_rates.json` in it.



## **G) Create a task named “creating\_forex\_rates\_table” - Hive operator and check it**

Now we need to create a HIVE table for forex data so we can query it.

```
from airflow.providers.apache.hive.operators.hive import HiveOperator
```

```
creating_forex_rates_table = HiveOperator(  
    task_id="creating_forex_rates_table",  
    hive_cli_conn_id="hive_conn",  
    hql=""
```

```

CREATE EXTERNAL TABLE IF NOT EXISTS forex_rates(
  base STRING,
  last_update DATE,
  eur DOUBLE,
  usd DOUBLE,
  nzd DOUBLE,
  gbp DOUBLE,
  jpy DOUBLE,
  cad DOUBLE
)
ROW FORMAT DELIMITED
FIELDS TERMINATED BY ','
STORED AS TEXTFILE
""""
)

```

In Airflow UI page navigate to Admin - > connections - > +

Conn Id \* : hive\_conn  
 Conn Type \* : Hive Server 2 Thrift  
 Host : hive-server  
 Login : hive  
 Password : hive  
 Port : 10000

Save

The screenshot shows the 'Edit Connection' form in the Airflow Admin UI. The form has the following fields and values:

- Conn Id \***: hive\_conn
- Conn Type \***: Hive Server 2 Thrift (dropdown menu)
- Description**: (empty text area)
- Host**: hive-server
- Schema**: (empty text field)
- Login**: hive
- Password**: (empty password field)
- Port**: 10000
- Extra**: (empty text area)

At the bottom left, there is a blue 'Save' button and a back arrow icon. A message below the Conn Type dropdown reads: 'Conn Type missing? Make sure you've installed the corresponding Airflow Provider Package.'

Test the task:

airflow tasks test forex\_data\_pipeline creating forex\_rates\_table 2023-01-01

```

[2023-08-02 10:26:00,421] {base.py:79} INFO - Using connection to: id: ***_conn. Host: ***-server, Port: 10000, Schema:
, Login: ***, Password: ***, extra: {}
[2023-08-02 10:26:00,421] {hive.py:155} INFO - Passing HiveConf: {'airflow.ctx.dag_email': 'admin@localhost.com', 'airfl
ow.ctx.dag_owner': 'airflow', 'airflow.ctx.dag_id': 'forex_data_pipeline', 'airflow.ctx.task_id': 'creating_forex_rates_
table', 'airflow.ctx.execution_date': '2023-01-01T00:00:00+00:00'}
[2023-08-02 10:26:00,422] {hive.py:247} INFO - *** -***conf airflow.ctx.dag_id=forex_data_pipeline -***conf airflow.ctx.
task_id=creating_forex_rates_table -***conf airflow.ctx.execution_date=2023-01-01T00:00:00+00:00 -***conf airflow.ctx.da
g_run_id= -***conf airflow.ctx.dag_owner=airflow -***conf airflow.ctx.dag_email=admin@localhost.com -***conf mapred.job.
name=Airflow HiveOperator task for 510f580614fb.forex_data_pipeline.creating_forex_rates_table.2023-01-01T00:00:00+00:00
-f /tmp/airflow_***op_0146a2no/tmp6hv24yq3
[2023-08-02 10:26:01,484] {hive.py:259} INFO - SLF4J: Class path contains multiple SLF4J bindings.
[2023-08-02 10:26:01,485] {hive.py:259} INFO - SLF4J: Found binding in [jar:file:/opt/***/lib/log4j-slf4j-impl-2.10.0.ja
r!/org/slf4j/impl/StaticLoggerBinder.class]
[2023-08-02 10:26:01,485] {hive.py:259} INFO - SLF4J: Found binding in [jar:file:/opt/hadoop/share/hadoop/common/lib/slf
4j-log4j12-1.7.25.jar!/org/slf4j/impl/StaticLoggerBinder.class]
[2023-08-02 10:26:01,485] {hive.py:259} INFO - SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an expla
nation.
[2023-08-02 10:26:01,497] {hive.py:259} INFO - SLF4J: Actual binding is of type [org.apache.logging.slf4j.Log4jLoggerFac
tory]
[2023-08-02 10:26:03,937] {hive.py:259} INFO - Hive Session ID = 90675f3c-06bd-4ed3-ad3f-a38c54df59bc
[2023-08-02 10:26:03,982] {hive.py:259} INFO -
[2023-08-02 10:26:03,983] {hive.py:259} INFO - Logging initialized using configuration in jar:file:/opt/***/lib/***/comm
on-3.1.2.jar!/***-log4j2.properties Async: true
[2023-08-02 10:26:06,310] {hive.py:259} INFO - Hive Session ID = 332028ad-56f4-4077-b839-d123fb1e7710
[2023-08-02 10:26:07,018] {hive.py:259} INFO - OK
[2023-08-02 10:26:07,018] {hive.py:259} INFO - Time taken: 0.666 seconds
[2023-08-02 10:26:07,666] {hive.py:259} INFO - OK
[2023-08-02 10:26:07,666] {hive.py:259} INFO - Time taken: 0.647 seconds
[2023-08-02 10:26:08,160] {taskinstance.py:1219} INFO - Marking task as SUCCESS. dag_id=forex_data_pipeline, task_id=cre
ating_forex_rates_table, execution_date=20230101T000000, start_date=20230802T102600, end_date=20230802T102600
airflow@510f580614fb:/$

```

In HUE we can see the HIVE table would have been created

Name	Size	User	Group	Permissions	Date
forex		root	supergroup	drwxr-xr-x	August 02, 2023 09:38 AM
.		airflow	supergroup	drwxr-xr-x	August 02, 2023 09:38 AM
forex_rates.json	245 bytes	airflow	supergroup	-rw-r--r--	August 02, 2023 09:38 AM

No we try to write query in HUE

SELECT \* FROM forex\_rates

Done. 0 results.

Query History

a few seconds ago SELECT \* FROM forex\_rates

We cannot see results as there is no data in table.

## H) Process the forex rates with Spark - Spark submit operator

```
from airflow.providers.apache.spark.operators.spark_submit import SparkSubmitOperator

forex_processing = SparkSubmitOperator(
    task_id="forex_processing",
    application="/opt/airflow/dags/scripts/forex_processing.py",
    conn_id="spark_conn",
    verbose=False
)
```

The below program basically creates a Spark session, read forex\_rates.json file and then do some pre processing on it and the finally insert it into forex\_rates table.

```
forex_processing.py

from os.path import expanduser, join, abspath

from pyspark.sql import SparkSession
from pyspark.sql.functions import from_json

warehouse_location = abspath('spark-warehouse')

# Initialize Spark Session

#warehouse_location?

spark = SparkSession \
    .builder \
    .appName("Forex processing") \
    .config("spark.sql.warehouse.dir", warehouse_location) \
    .enableHiveSupport() \
    .getOrCreate()

# Read the file forex_rates.json from the HDFS
df = spark.read.json('hdfs://namenode:9000/forex/forex_rates.json')

# Drop the duplicated rows based on the base and last_update columns
forex_rates = df.select('base', 'last_update', 'rates.eur', 'rates.usd', 'rates.cad', 'rates.gbp', 'rates.jpy',
'rates.nzd') \
    .dropDuplicates(['base', 'last_update']) \
    .fillna(0, subset=['EUR', 'USD', 'JPY', 'CAD', 'GBP', 'NZD'])

# Export the dataframe into the Hive table forex_rates
forex_rates.write.mode("append").insertInto("forex_rates")
```

In Airflow UI navigate to Admin - > Connections - > +

```
Conn Id * : spark_conn
Conn Type * : Spark
Host : spark://spark-master
Port : 7077
```

Save

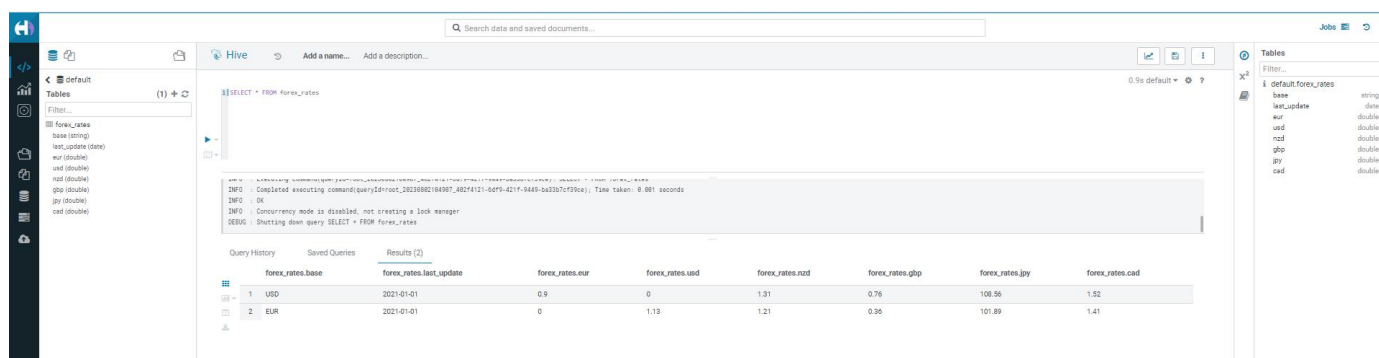
In cmd:

airflow tasks test forex\_data\_pipeline forex\_processing 2023-01-01

```
[2023-08-02 10:48:22,165] {spark_submit.py:523} INFO - 2023-08-02 10:48:22,165 INFO memory.MemoryStore: MemoryStore cleared
[2023-08-02 10:48:22,166] {spark_submit.py:523} INFO - 2023-08-02 10:48:22,165 INFO storage.BlockManager: BlockManager stopped
[2023-08-02 10:48:22,168] {spark_submit.py:523} INFO - 2023-08-02 10:48:22,168 INFO storage.BlockManagerMaster: BlockManagerMaster stopped
[2023-08-02 10:48:22,175] {spark_submit.py:523} INFO - 2023-08-02 10:48:22,175 INFO scheduler.OutputCommitCoordinator$OutputCommitCoordinatorEndpoint: OutputCommitCoordinator stopped!
[2023-08-02 10:48:22,187] {spark_submit.py:523} INFO - 2023-08-02 10:48:22,187 INFO spark.SparkContext: Successfully stopped SparkContext
[2023-08-02 10:48:22,187] {spark_submit.py:523} INFO - 2023-08-02 10:48:22,187 INFO util.ShutdownHookManager: Shutdown hook called
[2023-08-02 10:48:22,189] {spark_submit.py:523} INFO - 2023-08-02 10:48:22,188 INFO util.ShutdownHookManager: Deleting directory /tmp/spark-5cc3d773-520d-4ee0-93c0-30dc4f665e2a
[2023-08-02 10:48:22,192] {spark_submit.py:523} INFO - 2023-08-02 10:48:22,191 INFO util.ShutdownHookManager: Deleting directory /tmp/spark-7453cd12-98a8-4af4-b1cd-d80e27d4401a
[2023-08-02 10:48:22,194] {spark_submit.py:523} INFO - 2023-08-02 10:48:22,193 INFO util.ShutdownHookManager: Deleting directory /tmp/spark-5cc3d773-520d-4ee0-93c0-30dc4f665e2a/pyspark-b8bc2e69-1bee-4a70-9372-db92c41b9bc5
[2023-08-02 10:48:22,666] {taskinstance.py:1219} INFO - Marking task as SUCCESS. dag_id=forex_data_pipeline, task_id=forex_processing, execution_date=20230101T000000, start_date=20230802T104800, end_date=20230802T104822
airflow@510f580614fb:/$
```

Run the query in HUE, and we can see results as we have load data into the table

SELECT \* FROM forex\_rates



	forex_rates.base	forex_rates.last_update	forex_rates.eur	forex_rates.usd	forex_rates.rnd	forex_rates.gbp	forex_rates.jpy	forex_rates.cad
1	USD	2023-01-01	0.9	0	1.31	0.76	106.56	1.52
2	EUR	2023-01-01	0	1.19	1.21	0.36	101.89	1.41

## I) Send Email notifications - Email Operator

Now we need to configure our email provider so we can send email from our data pipeline by using our email address.

In the browser browse for <https://security.google.com/settings/security/apppasswords> and sign in to the gmail account.

App passwords

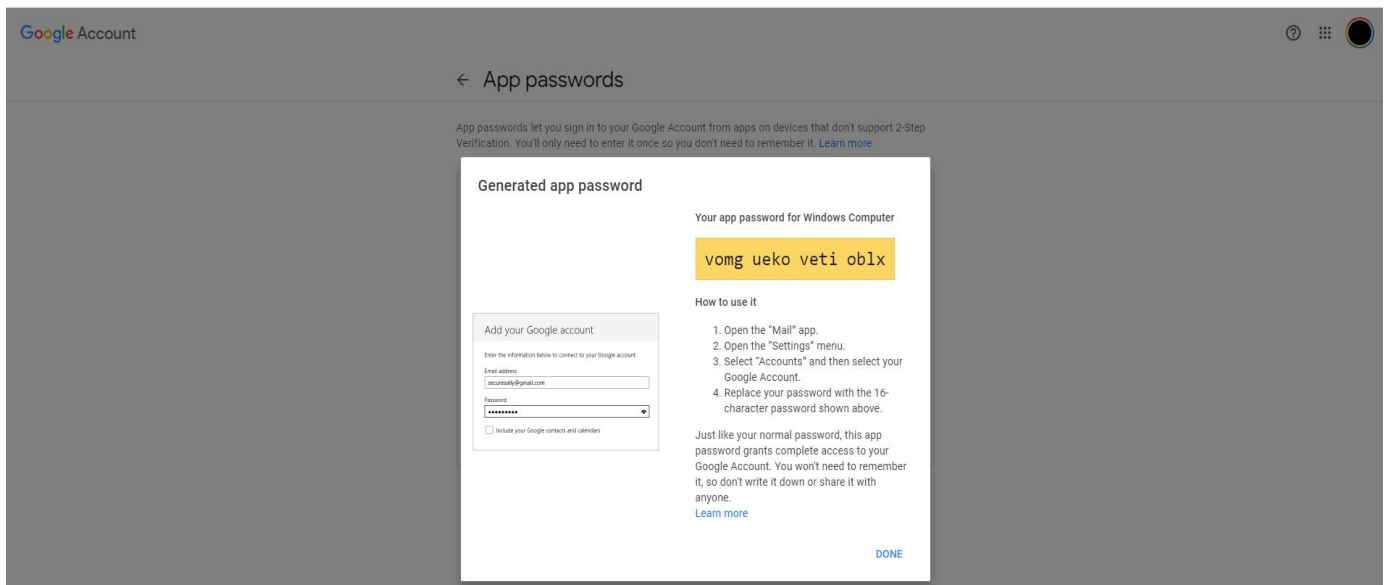
App passwords let you sign in to your Google Account from apps on devices that don't support 2-Step Verification.

Select App : Mail

Select device : Windows computer

Click on Generate

Click on Done



Configure SMTP in airflow.cfg located in “/mnt/airflow/”

```
[smtp]

# If you want airflow to send emails on retries, failure, and you want to use
# the airflow.utils.email.send_email_smtp function, you have to configure an
# smtp server here
smtp_host = smtp.gmail.com
smtp_starttls = True
smtp_ssl = False
# Example: smtp_user = airflow
smtp_user = naveenraju100@gmail.com
# Example: smtp_password = airflow
smtp_password = vomguekovetioblx
smtp_port = 587
smtp_mail_from = naveenraju100@gmail.com
smtp_timeout = 30
smtp_retry_limit = 5
```

Restart Docker container

In cmd

docker-compose restart airflow

```
from airflow.operators.email import EmailOperator

send_email_notification = EmailOperator(
    task_id="send_email_notification",
    to="nsreeramarajugovinda@hawk.iit.edu",
    subject="forex_data_pipeline",
    html_content="<h3>forex_data_pipeline</h3>"
)
```

In cmd : docker ps


# Copy the container id of airflow-section-3-airflow



```
docker exec -it 510f580614fb /bin/bash
```

```
airflow tasks test_forex_data_pipeline send_email_notification 2023-01-01
```

```
[2023-08-02 12:04:48,879] {taskinstance.py:1254} INFO - Exporting the following env vars:
AIRFLOW_CTX_DAG_EMAIL=admin@localhost.com
AIRFLOW_CTX_DAG_OWNER=airflow
AIRFLOW_CTX_DAG_ID=forex_data_pipeline
AIRFLOW_CTX_TASK_ID=send_email_notification
AIRFLOW_CTX_EXECUTION_DATE=2023-01-01T00:00:00+00:00
/usr/local/lib/python3.7/dist-packages/airflow/utils/email.py:102 PendingDeprecationWarning: Fetching SMTP credentials from configuration variables will be deprecated in a future release. Please set credentials using a connection instead.
[2023-08-02 12:04:48,883] {email.py:208} INFO - Email alerting: attempt 1
[2023-08-02 12:04:50,197] {email.py:220} INFO - Sent an alert email to ['nsreeramajugovinda@hawk.iit.edu']
[2023-08-02 12:04:51,866] {taskinstance.py:1219} INFO - Marking task as SUCCESS. dag_id=forex_data_pipeline, task_id=send_email_notification, execution_date=20230101T000000, start_date=20230802T120448, end_date=20230802T120451
airflow@510f580614fb:/$
```



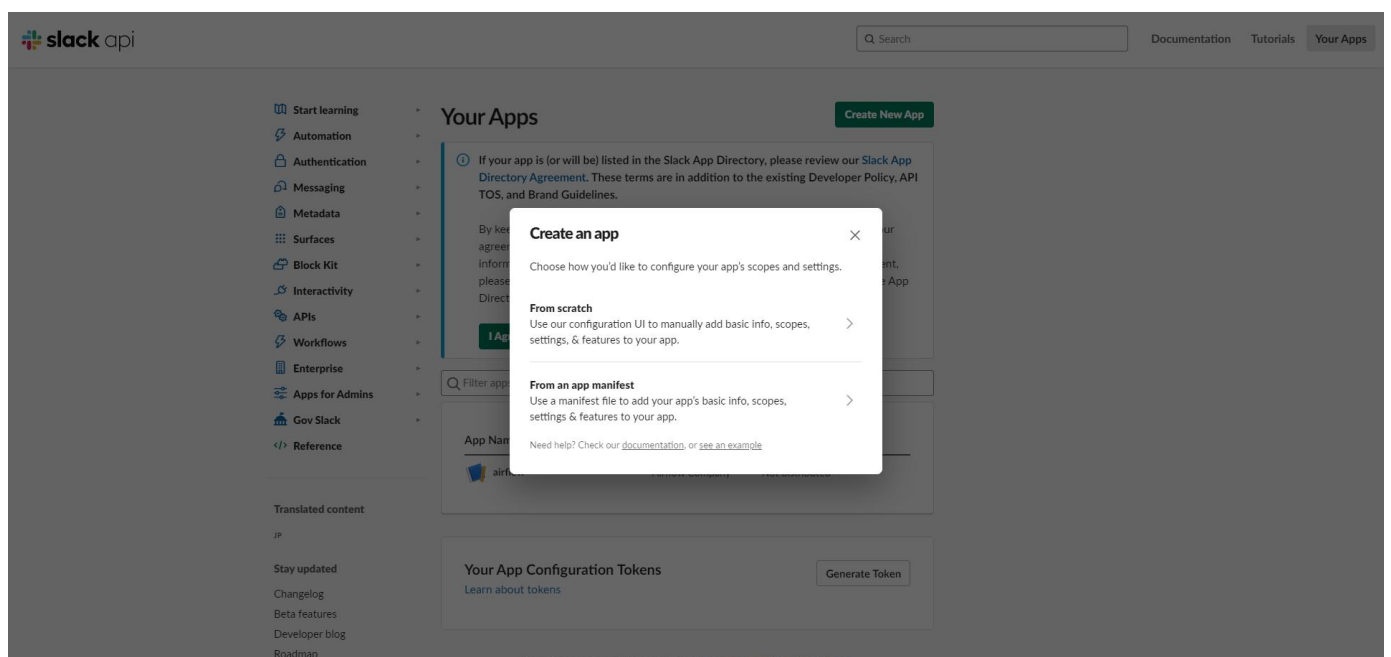
## J) Send Slack notifications - SlackWebHookOperator

Create a new a new Slack Workspace named “Airflow Company”

In the browser browse <https://api.slack.com/apps> and click “Create an App”

Select “From an app manifest”

Select “From Scratch”



## Name app & choose workspace


×

App Name

airflow

Don't worry - you'll be able to change this later.

Pick a workspace to develop your app in:

 Airflow Company

▼

Keep in mind that you can't change this app's workspace later. If you leave the workspace, you won't be able to manage any apps you've built for it. The workspace will control the app even if you leave the workspace.

[Sign into a different workspace](#)

By creating a Web API Application, you agree to the [Slack API Terms of Service](#).

CancelCreate App

Click on “Incoming Webhooks”

Incoming webhooks are a simple way to post messages from external sources into Slack.

airflow

Settings

Basic Information

Collaborators

Socket Mode

Install App

Manage Distribution

Features

App Home

Org Level Apps

Incoming Webhooks

Interactivity & Shortcuts

Slash Commands

Workflow Steps

OAuth & Permissions

Event Subscriptions

User ID Translation

App Manifest NEW

Beta Features

Submit to App Directory

Review & Submit

Give feedback

Slack ♥

Help

## Basic Information

### Building Apps for Slack

Create an app that's just for your workspace (or build one that can be used by any workspace) by following the steps below.

#### Add features and functionality

Choose and configure the tools you'll need to create your app (or review all our [documentation](#)).

**Building an internal app locally or behind a firewall?**

To receive your app's payloads over a WebSockets connection, enable [Socket Mode](#) for your app.

**Incoming Webhooks**

Post messages from external sources into Slack.

**Interactive Components**

Add components like buttons and select menus to your app's interface, and create an interactive experience for users.

**Slash Commands**

Allow users to perform app actions by typing commands in Slack.

**Event Subscriptions**

Make it easy for your app to respond to activity in Slack.

**Bots**

Allow users to interact with your app through channels and conversations. ✨

**Permissions**

Configure permissions to allow your app to interact with the Slack API.

Activate “Incoming Webhooks”  
Click on “Add New Webhook to Workspace”

airflow

Settings

Basic Information

Collaborators

Socket Mode

Install App

Manage Distribution

Features

App Home

Org Level Apps

Incoming Webhooks

Interactivity & Shortcuts

Slash Commands

Workflow Steps

OAuth & Permissions

Event Subscriptions

User ID Translation

App Manifest NEW

Beta Features

Submit to App Directory

Review & Submit

Give feedback

Slack ♥

Help

Contact

Policies

Our Blog

## Incoming Webhooks

### Activate Incoming Webhooks

On

Incoming webhooks are a simple way to post messages from external sources into Slack. They make use of normal HTTP requests with a JSON payload, which includes the message and a few other optional details. You can include [message attachments](#) to display richly-formatted messages.

Adding incoming webhooks requires a bot user. If your app doesn't have a bot user, we'll add one for you.

Each time your app is installed, a new Webhook URL will be generated.

If you deactivate incoming webhooks, new Webhook URLs will not be generated when your app is installed to your team. If you'd like to remove access to existing Webhook URLs, you will need to [Revoke All OAuth Tokens](#).

### Webhook URLs for Your Workspace

To dispatch messages with your webhook URL, send your [message](#) in JSON as the body of an [application/json](#) POST request.

Add this webhook to your workspace below to activate this curl example.

Sample curl request to post to a channel:

```
curl -X POST -H 'Content-type: application/json' --data '{"text":"Hello, World!"}'  
YOUR_WEBHOOK_URL_HERE
```

Webhook URL	Channel	Added By
No webhooks have been added yet.		

Add New Webhook to Workspace

#airflow is workspace created before



**airflow is requesting permission to access the  
Airflow Company Slack workspace**

**Where should airflow post?**

# airflow requires a channel to post to as an app

# airflow

Cancel

Allow

Copy the Webhook URL that is created

The screenshot shows the Slack 'Incoming Webhooks' settings page. On the left is a sidebar with navigation links: Settings (Basic Information, Collaborators, Socket Mode, Install App, Manage Distribution), Features (App Home, Org Level Apps, Incoming Webhooks, Interactivity & Shortcuts, Slash Commands, Workflow Steps, OAuth & Permissions, Event Subscriptions, User ID Translation, App Manifest, Beta Features), Submit to App Directory, and Give feedback. The main content area is titled 'Incoming Webhooks' and includes a toggle switch for 'Activate Incoming Webhooks' which is turned 'On'. Below this, there is explanatory text about webhooks and a sample curl request to post to a channel. At the bottom, there is a table listing existing webhooks.

Webhook URL	Channel	Added By
<a href="https://hooks.slack.com/services/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">https://hooks.slack.com/services/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</a> <span>Copy</span>	#airflow	NAVEEN RAJU Aug 2, 2023 <span>🗑️</span>

[Add New Webhook to Workspace](#)

<https://hooks.slack.com/services/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>

```
from airflow.providers.slack.operators.slack_webhook import SlackWebhookOperator
from airflow.hooks.base_hook import BaseHook

def _get_message() -> str:
    return "Hi from forex_data_pipeline"

send_slack_notification = SlackWebhookOperator(
    task_id="send_slack_notification",
    http_conn_id="slack_conn",
    message=_get_message(),
    channel="#airflow"
)
```



In Airflow UI navigate to Admin - > Connections - > +

```
Conn Id * : slack_conn
Conn Type * : HTTP
Host : https://hooks.slack.com/services
```



Password : [T05DN614PTJ/B05KHCBBUBZ/aSWGgf3bPHkjTRv7K1N5a2vI](#)

Save

 DAGs Security Browse Admin Docs 13:17 UTC 

Edit Connection

Conn Id \*

Conn Type \*   
Conn Type missing? Make sure you've installed the corresponding Airflow Provider Package.

Description

Host


Schema

Login

Password

Port

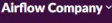
Extra



In cmd:

airflow tasks test forex\_data\_pipeline send\_slack\_notification 2023-01-01

```
use update your config:
[2023-08-02 13:15:26,271] {taskinstance.py:1254} INFO - Exporting the following env vars:
AIRFLOW_CTX_DAG_EMAIL=admin@localhost.com
AIRFLOW_CTX_DAG_OWNER=airflow
AIRFLOW_CTX_DAG_ID=forex_data_pipeline
AIRFLOW_CTX_TASK_ID=send_slack_notification
AIRFLOW_CTX_EXECUTION_DATE=2023-01-01T00:00:00+00:00
[2023-08-02 13:15:26,273] {base.py:79} INFO - Using connection to: id: slack_conn. Host: https://hooks.slack.com/service
s, Port: None, Schema: , Login: , Password: ***, extra: {}
[2023-08-02 13:15:26,274] {base.py:79} INFO - Using connection to: id: slack_conn. Host: https://hooks.slack.com/service
s, Port: None, Schema: , Login: , Password: ***, extra: {}
[2023-08-02 13:15:26,275] {http.py:140} INFO - Sending 'POST' to url: https://hooks.slack.com/services/**
[2023-08-02 13:15:26,736] {taskinstance.py:1219} INFO - Marking task as SUCCESS. dag_id=forex_data_pipeline, task_id=sen
d_slack_notification, execution_date=20230101T000000, start_date=20230802T122943, end_date=20230802T131526
airflow@510f580614fb:/$
```

 Upgrade Plan

Canvases  
Slack Connect  
Files  
Browse Slack

Channels

- # airflow
- # general
- # random
- + Add channels

Direct messages

- NAVEEN RAJU you
- + Add coworkers

Apps

- airflow
- airflow
- + Add apps

# airflow

+ Add a bookmark

Welcome to the # airflow channel

This channel is for everything #airflow. Hold meetings, share docs, and make decisions together with your team. [Edit description](#)

+ Add people

Today

NAVEEN RAJU 5:51 PM  
added an integration to this channel: airflow

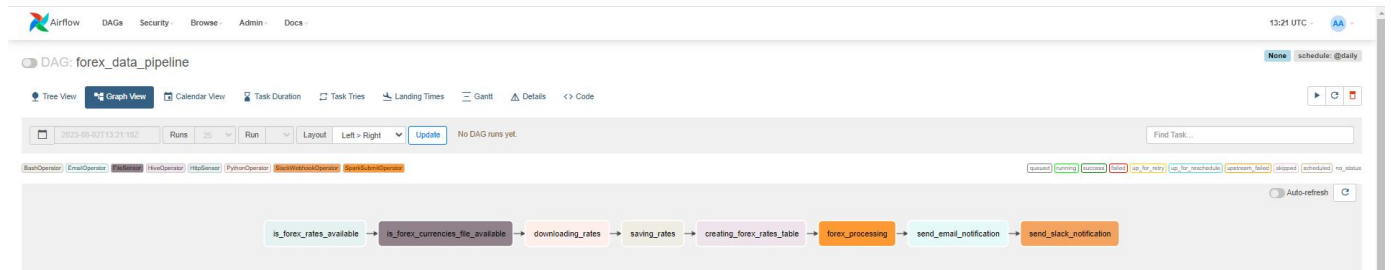
airflow APP 6:45 PM  
Hi from forex\_data\_pipeline

New

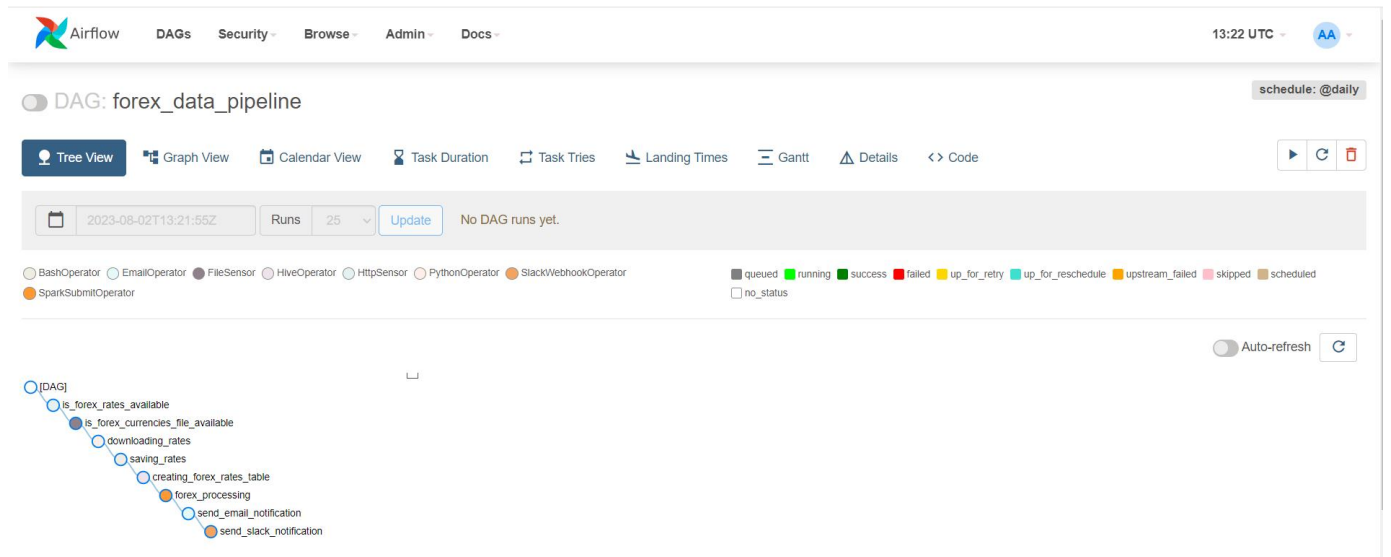
**K) Add dependency between tasks**

```
is_forex_rates_available >> is_forex_currencies_file_available >> downloading_rates >> saving_rates
saving_rates >> creating_forex_rates_table >> forex_processing
forex_processing >> send_email_notification >> send_slack_notification
```

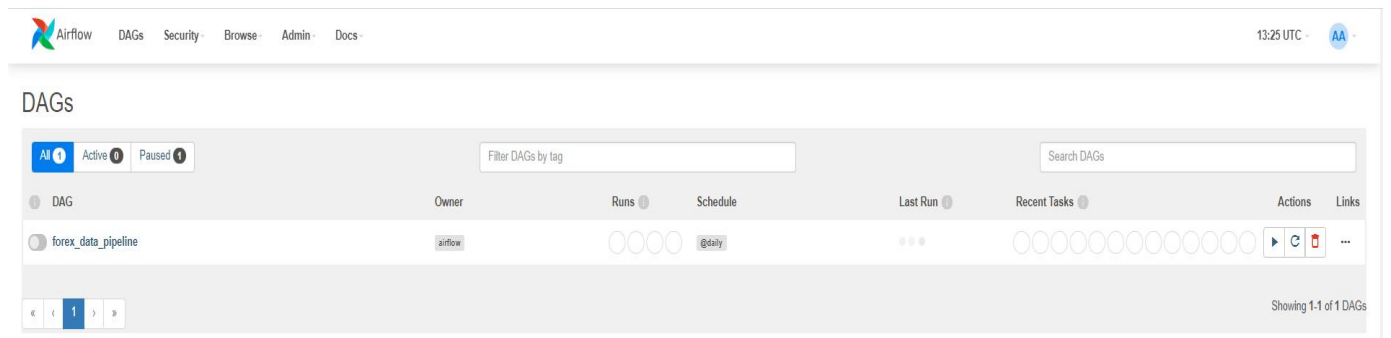
Graph view :




Tree view:



### **L) Trigger a DAG from airflow UI:**



Airflow

DAGs

Security

Browse

Admin

Docs

13:25 UTC

AA

DAG: forex\_data\_pipeline

Tree View

Graph View

Calendar View

Task Duration

Task Times

Landing Times

Gantt

Details

<> Code

2023-08-02T13:25:18Z

Runs

25

Update

No DAG runs yet.

BashOperator

EmailOperator

FileSensor

HiveOperator

HttpSensor

PythonOperator

StackWebhookOperator

SparkSubmitOperator

queued

running

success

failed

up\_for\_retry

up\_for\_reschedule

upstream\_failed

skipped

scheduled

no\_status

Auto-refresh

DAG

is\_forex\_rates\_available

is\_forex\_currency\_file\_available

downloading\_rates

extracting\_rates

creating\_forex\_rates\_table

forex\_processing

send\_email\_notification

send\_slack\_notification